

Playing the Minesweeper with Constraints

Raphaël Collet

Université catholique de Louvain,
B-1348 Louvain-la-Neuve, Belgium
`raph@info.ucl.ac.be`

Abstract. We present the design and implementation of a Minesweeper game, augmented with a digital assistant. The assistant uses constraint programming techniques to help the player, and is able to play the game by itself. It predicts safe moves, and gives probabilistic information when safe moves cannot be found.

1 Introduction

The Minesweeper game has been popular for several years now. Part of its popularity might come from its simplicity. A board represents a mine field, with mines hidden under the squares. The game consists in finding the mines without making them explode. You get new hints each time you uncover a non-mined square. Though, the simplicity does not make the game easy. The Minesweeper problem is hard: it has been proven NP-complete by Richard Kaye [1]. So simple techniques are not enough to solve it.

In this paper we show how the problem of finding safe moves can be modeled as a Constraint Satisfaction Problem (CSP). Techniques from the field of constraint programming can be used to program a digital assistant for a player. We applied several of them in a real application, the Oz Minesweeper [2]. This relatively small program demonstrates the power of the programming language Oz.

Contribution. The paper describes a diverting application, that applies various techniques from constraint programming to implement a digital assistant. The contribution is mainly educational. We model a simple problem, and show the key ideas underlying the application's implementation.

History. The Oz Minesweeper is born in spring 1998. It started as a student work, in a course on constraint programming. The goal was to study a programming language called Oz, and give a presentation about it. To make the presentation attractive, I showed an example of a CSP in the “real world”, namely the Minesweeper game. I had hacked a small solver that was playing the game.

Later I rewrote it as a demonstration program, and gave it a graphical user interface. A simple inference engine based on propagation was provided. It already impressed quite a lot of visitors. The next step was a solver, which was basically making the inference engine complete. I wrote several implementations

of it, notably by hacking a special search engine. I eventually found a way to compute mine probabilities. I rewrote everything from scratch. The hacked special search engine went to the trash can.

The last step happened last year. I understood the issue of symmetries in the problem, and designed an improved solver that eliminates them. A better propagation-based inference engine was designed while implementing the solver. I reworked a bit the implementation, and integrated the inference engines in a proper way. I finally improved the user interface the week before submitting this paper.

Paper Organization. Section 2 recalls the rules of the game, and proposes a simple mathematical model for it. Section 3 investigates how constraint programming techniques can be applied in order to solve the problem with reasonable efficiency. Section 4 then gives an overview of the implementation of the Oz Minesweeper. Section 5 evaluates and quickly compares our work to other similar products.

2 The Game as a Constraint Satisfaction Problem

Let us recall the rules of the game. A mine field is given to the player as a rectangular board. Each square on the board may hide at most one mine. The total number of mines is known by the player. A move consists in uncovering a square. If the square holds a mine, the mine explodes and the game is over. Otherwise, a number in the square indicates how many mines are held in the surrounding squares, which are the adjacent squares in the eight directions north, north-east, east, south-east, south, south-west, west, and north-west. The goal of the game is to uncover all the squares that do not hold a mine.




Fig. 1. An example of a board with 20 mines

Figure 1 shows an example of a board that contains 20 mines. We identify each square by its coordinates (*row, column*). The squares (1,1), (1,2), (1,3), (1,4), (2,1), (2,2), (2,3), (2,4), (3,4), and (4,4) have already been played, and have no mine in their respective surrounding squares. The squares (1,5), (3,1), (3,2), (3,3), (4,3), and (4,5) have been played, too, and are surrounded by one mined square each. The squares (2,5), (3,5), (5,4), and (5,5) each have two mines in their neighborhood, while the square (5,3) has three mines around it. In this example, the player might deduce from (3,3) that (4,2) is mined, and by (3,2) that (4,1) is a safe move.

Model. Finding safe moves on the board consists in solving the problem defined by those numbers in the squares. The unknown of the problem is the positions of the mines. We model this as a binary matrix that represents the mine field, with one entry per square. The value 1 means that the corresponding square is mined, while 0 means a safe square.

$$\begin{pmatrix} x_{11} & x_{12} & \cdots & x_{1n} \\ x_{21} & x_{22} & & x_{2n} \\ \vdots & & \ddots & \vdots \\ x_{m1} & x_{m2} & \cdots & x_{mn} \end{pmatrix}$$

By convention, x_{ij} always denotes the matrix entry corresponding to the square at position (i, j) . The problem can be written as linear equations over the x_{ij} 's. In the example, we have 20 mines (first equation below), and the played squares are not mined (second equation below). The other equations are given by the numbers in the squares. The corresponding square coordinates are given on the left of each equation.

$$\begin{aligned} \sum_{i,j \in \{1, \dots, 10\}} x_{ij} &= 20 \\ x_{11} &= x_{12} = x_{13} = \cdots = 0 \\ (1, 1) \quad x_{12} + x_{21} + x_{22} &= 0 \\ (1, 2) \quad x_{11} + x_{13} + x_{21} + x_{22} + x_{23} &= 0 \\ &\dots \\ (1, 5) \quad x_{14} + x_{16} + x_{24} + x_{25} + x_{26} &= 1 \\ &\dots \\ (2, 5) \quad x_{14} + x_{15} + x_{16} + x_{24} + x_{26} + x_{34} + x_{35} + x_{36} &= 2 \\ &\dots \end{aligned}$$

This binary model of the game defines a CSP, which we can solve to find hints for the player's next move. If we have all the solutions of the problem, we can look at what is common to all those solutions. For instance, if all solutions give $x_{41} = 0$, we know that the square at position (4,1) is a safe move.

But we even go further than this. Assuming that all those solutions have the same probability, we can compute the expected solution, i.e., the mean of all solutions. This gives us a probability for each square to be mined. In case no safe move can be found, the player might use this information to choose her next move.

3 Propagation, Search, and Probabilities

We now present specific information related to the implementation of the inference engines. Each of them provides a way to solve the CSP defined by the current state of the game. Sections 3.1 and 3.2 shows two implementation based on constraint propagation only. Sections 3.3 and 3.4 presents two solvers, and explains how their results are used to compute mine probabilities.

3.1 Simple Propagators

The simplest inference engine uses the binary model of the Minesweeper game, and posts the propagators that trivially implement the constraints of the model. We illustrate this with the example shown in Fig. 1. A quick sketch of the CSP is given at the end of Sect. 2. All those constraints can be implemented with the Oz propagator `FD.sum`, taking a list of `FD` variables with domain `0#1`. For instance, the propagator for (2,5) is created by a statement like

```
{FD.sum [X14 X15 X16 X24 X26 X34 X35 X36] '=: ' 2}
```

Let us examine the effect of those propagators. For the sake of simplicity, we assume that the “zero” constraints like (1,1) have been propagated, and we simplify the remaining constraints using the known values. The constraints are

$$\begin{array}{ll}
 (1, 5) & x_{16} + x_{26} = 1 \\
 (2, 5) & x_{16} + x_{26} + x_{36} = 2 \\
 (3, 1) & x_{41} + x_{42} = 1 \\
 (3, 2) & x_{41} + x_{42} = 1 \\
 (3, 3) & x_{42} = 1 \\
 (3, 5) & x_{26} + x_{36} + x_{46} = 2 \\
 (4, 3) & x_{42} + x_{52} = 1 \\
 (4, 5) & x_{36} + x_{46} + x_{56} = 1 \\
 (5, 3) & x_{42} + x_{52} + x_{62} + x_{63} + x_{64} = 3 \\
 (5, 4) & x_{63} + x_{64} + x_{65} = 2 \\
 (5, 5) & x_{46} + x_{56} + x_{64} + x_{65} + x_{66} = 2
 \end{array}$$

The propagator for (3,3) immediately infers $x_{42} = 1$, which means that we have found the position of a mine. This information allows propagators (3,1) and (3,2) to infer $x_{41} = 0$, while the propagator (4,3) infers $x_{52} = 0$.

The remaining propagators cannot infer new constraints, and thus wait for more information to come. Still, more information can be deduced from those constraints. But the propagators that we have considered here cannot do it, because they share too few information with each other. For instance, propagators (1,5) and (2,5) could infer $x_{36} = 1$ if they were sharing $x_{16} + x_{26} = 1$ as a basic constraint. This insight leads us to an improvement in the propagation of the constraints.

3.2 The Set Propagators

We now show propagators that infer information about sets of squares, hence the name “set” propagators. We continue with the example shown in Fig. 1. Let us assume that the simple propagators have determined the variables as explained above. We consider the remaining constraints

$$\begin{array}{ll}
 (1, 5) & x_{16} + x_{26} = 1 \\
 (2, 5) & x_{16} + x_{26} + x_{36} = 2 \\
 (3, 5) & x_{26} + x_{36} + x_{46} = 2 \\
 (4, 5) & x_{36} + x_{46} + x_{56} = 1
 \end{array}
 \qquad
 \begin{array}{ll}
 (5, 3) & x_{42} + x_{52} + x_{62} + x_{63} + x_{64} = 3 \\
 (5, 4) & x_{63} + x_{64} + x_{65} = 2 \\
 (5, 5) & x_{46} + x_{56} + x_{64} + x_{65} + x_{66} = 2
 \end{array}$$

Remember that the weakness of the simple propagators was coming from their inability to share information about subterms like $x_{16} + x_{26}$. Consider for instance constraint (2,5). The improved implementation of this constraint will actually create as many propagators as partitions of the set of indices $\{16, 26, 36\}$.

For each subset I of indices, we consider the “set” variable x_I defined by

$$x_I = \sum_{i \in I} x_i \quad (0 \leq x_I \leq |I|).$$

The definition of x_I can be implemented by a simple propagator over finite integers. We can now express the constraint (2,5) as follows. For each partition $P = \{I_1, \dots, I_k\}$ of the indices, we create one propagator for the constraint

$$x_{I_1} + \dots + x_{I_n} = 2,$$

which is logically equivalent to (2,5). We thus have propagators for the following equations. Note that (2,5)(a) has the same effect as the simple propagator for (2,5).

$$\begin{array}{ll}
 (2, 5)(a) & x_{\{16\}} + x_{\{26\}} + x_{\{36\}} = 2 \\
 (2, 5)(b) & x_{\{16\}} + x_{\{26, 36\}} = 2 \\
 (2, 5)(c) & x_{\{26\}} + x_{\{16, 36\}} = 2 \\
 (2, 5)(d) & x_{\{36\}} + x_{\{16, 26\}} = 2 \\
 (2, 5)(e) & x_{\{16, 26, 36\}} = 2
 \end{array}$$

Let us observe the effect of those propagators in the example. One of the propagators for (1,5) infers $x_{\{16, 26\}} = 1$, which makes (2,5)(d) infer $x_{\{36\}} = 1$, giving $x_{36} = 1$. The simple propagator (4,5) then infers $x_{46} = x_{56} = 0$. The propagation of (3,5) and (1,5) then gives $x_{26} = 1$ and $x_{16} = 0$.

3.3 A Binary Solver

As we said in Sect. 2, useful information can be deduced from the set of solutions of the Minesweeper problem. The issue is, there usually are *many* solutions. Consider the board in Fig. 2, which contains 20 mines. Four squares have been played. It defines the following CSP.

$$\begin{array}{ll}
 \sum_{i,j \in \{1, \dots, 10\}} x_{ij} = 20 \\
 x_{11} = x_{21} = x_{31} = x_{32} = 0 \\
 (1, 1) & x_{12} + x_{21} + x_{22} = 1 \\
 (2, 1) & x_{11} + x_{12} + x_{22} + x_{31} + x_{32} = 1 \\
 (3, 1) & x_{21} + x_{22} + x_{32} + x_{41} + x_{42} = 1 \\
 (3, 2) & x_{21} + x_{22} + x_{23} + x_{31} + x_{33} + x_{41} + x_{42} + x_{43} = 3
 \end{array}$$




Fig. 2. An example for search

Table 1. Solutions of the restricted binary problem

solution	s_1	s_2	s_3	s_4	s_5	s_6	s_7	s_8	s_9
x_{12}	0	0	0	1	1	1	1	1	1
x_{22}	1	1	1	0	0	0	0	0	0
x_{23}	0	1	1	0	1	1	0	1	1
x_{33}	1	0	1	1	0	1	1	0	1
x_{41}	0	0	0	0	0	0	1	1	1
x_{42}	0	0	0	1	1	1	0	0	0
x_{43}	1	1	0	1	1	0	1	1	0
class size	$\binom{89}{17}$	$\binom{89}{17}$	$\binom{89}{17}$	$\binom{89}{16}$	$\binom{89}{16}$	$\binom{89}{16}$	$\binom{89}{16}$	$\binom{89}{16}$	$\binom{89}{16}$

This problem has 3.333×10^{18} solutions. Computing all solutions is simply impossible, except for very small boards.

Though, that issue can be addressed. We simply restrict the problem to some of its variables. Each solution of the restricted problem defines a class of solutions of the full problem¹. We consider the variables given by the squares neighboring the already played squares. In the example, this gives x_{12} , x_{22} , x_{23} , x_{33} , x_{41} , x_{42} , and x_{43} . The remaining unknowns can be determined by other simple means. The solutions of the restricted problem are given in Table 1. If we consider the solution s_1 in the table, there remains 89 unknowns, out of which 17 must be mined. The number of ways to choose 17 elements out of 89 is given by the binomial $N_1 = \binom{89}{17}$. This is the size of the class of solutions defined by s_1 . The same argument applies for all classes.

¹ The word “class” is used with the meaning of “subset” here. The subset we consider is actually an equivalence class in the set of solutions.

The problem clearly has $N = N_1 + N_2 + \cdots + N_9 \simeq 3.333 \times 10^{18}$ solutions. Let X_{ij} denote how many solutions satisfy $x_{ij} = 1$. The probability that $x_{ij} = 1$ is simply given by X_{ij}/N . As a first example, take square (1,2). As $x_{12} = 0$ in s_1, s_2 , and s_3 , we have $X_{12} = N_4 + N_5 + \cdots + N_9 \simeq 1.059 \times 10^{18}$. So the probability that $x_{12} = 1$ is $X_{12}/N \simeq 0.318$. Now take square (6,9). In the class s_1 , the number of solutions satisfying $x_{69} = 1$ is $\binom{88}{16} = \frac{17}{89} \binom{89}{17}$. Summing up those numbers for all classes of solutions, we have $X_{69} \simeq 6.247 \times 10^{17}$, which gives a mine probability of 0.187.

3.4 The Set Solver

The binary solver still computes too many solutions. In the example, one can see that the problem has *symmetries*. For instance, each permutation of the values of x_{23}, x_{33}, x_{43} in one solution leads to another solution. This symmetry comes from the fact that those three variables are constrained by $x_{23} + x_{33} + x_{43} = 2$ only.

The improved solver reformulates the CSP in terms of the set variables x_I in order to eliminate those symmetries. Taking all equations that define the binary problem, it computes a partition of the variable's indices. Every subset I in the partition is such that, for each equation $x_J = k$ in the problem, $I \cap J = I$ or \emptyset . The subsets are chosen to be *maximal*, so that symmetries are eliminated.

Let us reformulate the CSP of the example in Fig. 2, which gives

$$\begin{aligned} \sum_{I \in P} x_I &= 20 \\ x_{\{11\}} &= x_{\{21\}} = x_{\{31\}} = x_{\{32\}} = 0 \\ (1, 1) \quad x_{\{12\}} + x_{\{21\}} + x_{\{22\}} &= 1 \\ (2, 1) \quad x_{\{11\}} + x_{\{12\}} + x_{\{22\}} + x_{\{31\}} + x_{\{32\}} &= 1 \\ (3, 1) \quad x_{\{21\}} + x_{\{22\}} + x_{\{32\}} + x_{\{41,42\}} &= 1 \\ (3, 2) \quad x_{\{21\}} + x_{\{22\}} + x_{\{31\}} + x_{\{41,42\}} + x_{\{23,33,43\}} &= 3 \end{aligned}$$

The indices have been partitioned into

$$P = \{\{11\}, \{12\}, \{21\}, \{22\}, \{31\}, \{32\}, \{41, 42\}, \{23, 33, 43\}, R\},$$

where R contains the remaining indices. This problem has two solutions, shown in Table 2. Each class of solutions is equivalent to the Cartesian product of the possible combinations for the set variables of the reformulated problem. Each valuation $x_I = k$ has $\binom{n}{k}$ solutions, where $n = |I|$. Therefore the size of each class is given by a product of binomials. The computation of the probabilities is similar to what the binary solver does. For instance, the probability that $x_{41} = 1$ is $(\frac{0}{2}N_1 + \frac{1}{2}N_2)/N \simeq 0.159$.

The efficiency is typically one order of magnitude faster compared to the binary solver. Let us illustrate this with an example. Figure 3 shows a snapshot of the application's window. The squares containing a mine have been marked with a black disk. The probabilities are drawn as filled rectangles in the squares. The more a rectangle is filled, the greater the mine probability. A precise value

Table 2. Solutions of the reformulated problem

solution	s_1	s_2	
$x_{\{12\}}$	0	1	$N_1 = \binom{1}{0} \binom{1}{1} \binom{2}{0} \binom{3}{2} \binom{89}{17}$
$x_{\{22\}}$	1	0	
$x_{\{41,42\}}$	0	1	$N_2 = \binom{1}{1} \binom{1}{0} \binom{2}{1} \binom{3}{2} \binom{89}{16}$
$x_{\{23,33,43\}}$	2	2	
x_R	17	16	$N = N_1 + N_2$
class size	N_1	N_2	

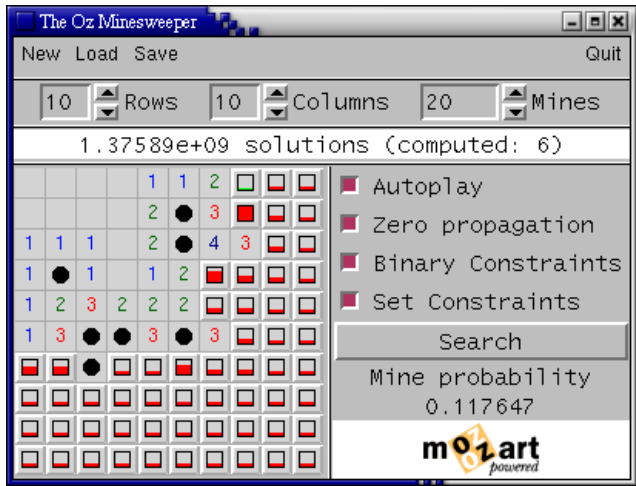


Fig. 3. A snapshot showing probabilistic information

of a probability is shown in the bottom right of the window when the player moves her mouse cursor over a given square. The set solver has computed 6 solutions to find the probabilities, while the binary solver would compute 246 solutions for the same problem!

4 Implementation

The general architecture of the Oz Minesweeper is depicted in Fig. 4. Boxes refer to concurrent agents (active objects), while “Symbolic field” and “Symbolic constraints” are simply shared data. Arrows from data to agents (resp. from agents to data) correspond to *ask* (resp. *tell*) operations. Arrows between agents represent messages or procedure calls. The removal of the components in the dashed box gives an implementation of the game without digital assistance.

4.1 The Core Components

The central point in the application is the *symbolic field*, which simply reflects the information known about the mine field. The symbolic field is a tuple whose elements correspond to the board squares. An element can be either `safe(κ)` or `mine(x)`. The value `safe(κ)` means that the corresponding square is not mined, and κ gives the number of mines in the surrounding squares. Note that κ can be unbound, if the square is known to be safe, but has not been played yet. The value `mine(x)` means that the square is mined, and x is bound to `exploded` if the mine has exploded, i.e., the game is over.

The *user interface* updates the board by threads that synchronize on the symbolic field. For instance, if an entry in the symbolic field is `safe(κ)` and κ is unbound, the square is marked with a dash “-”. This shows the user that this square is safe. As soon as κ is determined, its value is shown in the square, which becomes inactive. When the user clicks on a square, the user interface calls the *game* agent to play that square. The game automatically tells the result in the symbolic field, which wakes up the thread that updates the square.

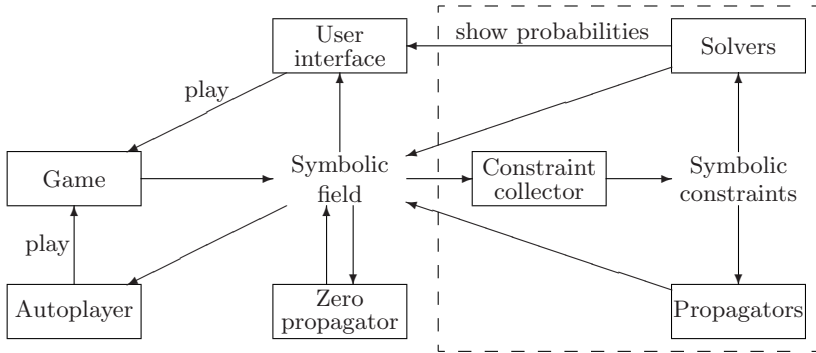


Fig. 4. Dataflow diagram of the Oz Minesweeper

4.2 The Zero Propagator and Autoplayer

The *zero propagator* simply asks and tells information in the symbolic field. If a square has no mines around it, which correspond to value `safe(0)` in the symbolic field, the surrounding squares are told to be safe. The code of the propagator is shown below. The symbolic field appears as the tuple `SymField`. The procedure `waitEnabled` blocks until the user enables the propagator. The same mechanism is used by all inference engines, and allows to user to experiment with them. The call to function `BoxI` returns the coordinates of all the squares in a box around square `I`.

```

for I in 1..{Width SymField} do
  thread
    case SymField.I of safe(0) then
      {WaitEnabled}
      for J in {BoxI I 1} do SymField.J = safe(_) end
    else skip end
  end
end

```

The *autoplayer* works in a similar way. When enabled, it plays all the squares known to be safe in the symbolic field. So the user can let the various inference engines discover safe moves, and decide whether they should be played automatically.

4.3 The Constraint Inference Engines

The *constraint collector* incrementally builds the *symbolic constraints*, a list of the constraints that appear implicitly in the symbolic field. The inference engines using constraint programming simply read this list to get the constraints of the current problem. A constraint in the list has the form $\text{sum}(\text{Is } \kappa)$, where Is is a list of square coordinates, and κ is a nonnegative integer. Its semantics is the equation $\sum_{i \in \text{Is}} x_i = \kappa$. All the constraints in the Minesweeper problem can be written in this way.

Propagators. Both the simple and set propagators read the symbolic constraints and progressively post propagators as explained in Sect. 3. Those propagators are posted over binary constrained variables, that correspond to the x_{ij} 's in the model. Whenever such a variable is determined, the information is automatically told in the symbolic field with a statement like

```

thread
  SymField.I = if X.I==0 then safe(_) else mine(_) end
end

```

Recall that the set propagator for the equation $\sum_{i \in I} x_i = k$ reformulates it as $x_{I_1} + \dots + x_{I_n} = k$, for every partition $\{I_1, \dots, I_n\}$ of I . If I has 8 elements (the typical case in the Minesweeper), this gives 255 set variables, and 4140 equations! The implementation optimizes this simple scheme. First, the equation is simplified by subtracting the known x_i 's. Second, the set variables are created lazily, and memoized for sharing between propagators. When a set variable x_I is created, a propagator is posted for $x_I = \sum_{i \in I} x_i$.

Solvers. A search with a solver is triggered by pushing a button in the user interface. The solver first takes the known part of the symbolic constraints list, and solves the problem given by those constraints. In the case of the set solver, it implies to first compute the optimal partition of the indices, to reformulate the constraints in terms of the set variables, and to solve the reformulated problem. If new safe moves or mine positions are found, they are told to the symbolic field. Otherwise, the mine probabilities are shown on the board.

5 Evaluation and Related Work

The Oz Minesweeper has been entirely written in Mozart [3], and is about 1000 lines of code. The digital assistant is capable to find all the safe moves in a given situation. The set propagator proved to be effective at this task, it usually finds most of them. The solver rarely finds new moves, and provides mine probabilities instead. It leaves the player with the toughest decision, involving a cost-benefit strategy. An interesting observation we have made is that the proportion of mined squares should be around 20% to make the game interesting. A proportion less than 20% makes the problem too easy, while more than 20% quickly makes the game unplayable.

We have not explored the problem of choosing a square to play when all you know is the mine probabilities. Playing the square with the lowest mine probability is a safe and conservative move. But we observed that such moves do not often bring much information to go further. It seems that taking a risk can be worth the candle. Our implementation is flexible enough to implement strategies on top of the existing solvers, which would provide a complete digital player.

We have found only one other application that solves the Minesweeper problem and computes the mine probabilities, called *Truffle-Swine Keeper* [4]. It seems efficient, but we have found the interaction with the solver not as practical as the Oz Minesweeper.

Other Techniques. Is constraint programming really a good choice for solving this problem? The Minesweeper problem is completely defined by linear equations. So one might think that integer programming could be a better choice. To my current understanding, integer programming can be applied successfully for some parts, but not all of them. We can use integer programming for checking a board square, for instance. Given a problem P and a variable x_{ij} , we check whether $P \wedge x_{ij} = 0$ is solvable. In case it is not, we can infer $x_{ij} = 1$ in our case. But I don't see how to use it for computing mine probabilities. The latter is a result about *all* solutions of a problem, while integer programming is oriented toward finding *one* solution.

6 Conclusion

We have designed and implemented a Minesweeper application with a digital assistant. The latter is based on a simple mathematical model of the Minesweeper game, and various techniques coming from the field of constraint programming. It proved to be effective, and is capable to infer every logical consequence of the problem to solve. It computes mine probabilities without computational burden.

The simplicity and efficiency of our application relies on the language Oz and the platform Mozart. The dataflow concurrency, symbolic data, and constraint system make the application's architecture modular and elegant.

References

1. Kaye, R.: Minesweeper is NP-complete. *Mathematical Intelligencer* (2000) See also <http://web.mat.bham.ac.uk/R.W.Kaye/minesw/ordmsw.htm> (08/26/2004).
2. Collet, R.: The Oz Minesweeper (2004) Program available at <http://www.info.ucl.ac.be/~raph/minesweeper/> (08/26/2004).
3. Mozart Consortium (DFKI, SICS, UCL, UdS): The Mozart programming system (Oz 3) (1999) Available at <http://www.mozart-oz.org>.
4. Kopp, H.: Truffle-Swine Keeper (2001) Program available at <http://people.freenet.de/hskopp/swinekeeper.html> (08/26/2004).